

Inverses

Reading: Strang 2.5

Learning objective: Understand when a product of matrices is invertible. Be familiar with the “big list” of conditions equivalent to invertibility.

Review

Definition: A square matrix A is **invertible** if and only if there is a matrix B such that

$$AB = I \quad \text{and} \quad BA = I$$

Here I is the identity matrix of the same size as A .

An n -by- n matrix A is invertible if and only if Gaussian elimination produces n pivots.

An n -by- n matrix A is invertible if and only if Gaussian elimination produces n pivots.

1) The “easy” direction: if A is invertible then Gaussian elimination produces n pivots.

because $A\vec{x} = \vec{0}$ has a unique solution.

2) If elimination produces n pivots, then A has a right inverse.

because $A\vec{x} = \vec{b}$ has a solution for any \vec{b} .

3) If elimination produces n pivots, then A has a left inverse.

because Gauss-Jordan produces the identity matrix and is done by left multiplication by elem. matrices.

Computing the Inverse

We can compute the inverse of A by solving $A\vec{x} = \vec{e}_i$, where \vec{e}_i is the i^{th} column of the identity matrix, for every i .

These solutions form the columns of A^{-1} .

We can organize this computation by forming the super-augmented matrix.

$$\begin{bmatrix} 1 & 1 & 1 & 1 & 0 & 0 \\ 1 & 2 & 4 & 0 & 1 & 0 \\ 1 & 3 & 9 & 0 & 0 & 1 \end{bmatrix}$$

Right vs Left

$$\begin{bmatrix} 1 & 1 & 1 & 1 & 0 & 0 \\ 1 & 2 & 4 & 0 & 1 & 0 \\ 1 & 3 & 9 & 0 & 0 & 1 \end{bmatrix}$$

explains gaussian jordan elimination

This computation is solving for a right inverse

$$AX = I$$

But we can also see why it gives a left inverse.

The steps of Gauss-Jordan elimination do

$$AX = I \rightarrow E_1 AX = E_1 \rightarrow E_k \cdots E_2 E_1 AX = E_k \cdots E_2 E_1$$

If $E_k \cdots E_2 E_1 A = I$ then $X = E_k \cdots E_2 E_1$.

Homogeneous Equations

A homogeneous system of equations is one where the right hand side vector is the zero vector $\vec{0}$.

In this case, the right hand side remains $\vec{0}$ throughout Gaussian elimination.

Theorem: A square matrix A is invertible if and only if

$$A\vec{x} = \vec{0}$$

has a unique solution.

Reason: $A\vec{x} = \vec{0}$ has a unique solution if and only if A has a pivot in every column.

Question

Suppose that A and B are invertible.

Is AB invertible?

Question

Suppose that B is singular.

Is AB invertible for some A ?

Left Inverse Suffices

Normally to show the square matrix A is invertible, we have to find B such that

$$AB = I \quad \text{and} \quad BA = I$$

Theorem: If $BA = I$ for square matrices A, B then A is invertible.

Proof: We show the contrapositive. If A is singular then there is a vector $\vec{u} \neq \vec{0}$ with $A\vec{u} = \vec{0}$.

Then $BA\vec{u} = B\vec{0} = \vec{0}$, thus BA is singular and cannot be the identity matrix.

Right Inverse Suffices

Suppose that A has a right inverse. There is a matrix B such that $AB = I$.

This means B has a **left inverse**, and so B is invertible.

As B is invertible, it also has a right inverse, which must equal its left inverse A (left inverse = right inverse).

This means $BA = I$ and A is invertible.

The Big List

Let A be a square matrix. The following are equivalent:

- A is **invertible**.
- Gaussian elimination produces a full set of pivots.
- $A\vec{x} = \vec{0}$ has a **unique** solution.
- A has a left inverse.
- A has a right inverse.
- The reduced row echelon form of A is the identity matrix.

LU Decomposition

Reading: Strang 2.6

Learning objective: Be able to compute the LU decomposition of a matrix and understand its application to solving systems of linear equations.

Example

$$\begin{bmatrix} 1 & 1 & 1 & 1 \\ 1 & 2 & 3 & 4 \\ 1 & 3 & 6 & 10 \\ 1 & 4 & 10 & 20 \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 1 & 1 & 0 & 0 \\ 1 & 2 & 1 & 0 \\ 1 & 3 & 3 & 1 \end{bmatrix} \begin{bmatrix} 1 & 1 & 1 & 1 \\ 0 & 1 & 2 & 3 \\ 0 & 0 & 1 & 3 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

We will find this beautiful LU decomposition.

We start with Gaussian elimination and keep track of the elementary matrices implementing the row operations.

Example

$$\begin{bmatrix} 1 & 1 & 1 & 1 \\ 1 & 2 & 3 & 4 \\ 1 & 3 & 6 & 10 \\ 1 & 4 & 10 & 20 \end{bmatrix} \xrightarrow{\begin{array}{l} R'_2 = R_2 - R_1 \\ R'_3 = R_3 - R_1 \\ R'_4 = R_4 - R_1 \end{array}} \begin{bmatrix} 1 & 1 & 1 & 1 \\ 0 & 1 & 2 & 3 \\ 0 & 2 & 5 & 9 \\ 0 & 3 & 9 & 19 \end{bmatrix}$$

Multiplying the corresponding elementary matrices together we have

$$\begin{bmatrix} 1 & 0 & 0 & 0 \\ -1 & 1 & 0 & 0 \\ -1 & 0 & 1 & 0 \\ -1 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} 1 & 1 & 1 & 1 \\ 1 & 2 & 3 & 4 \\ 1 & 3 & 6 & 10 \\ 1 & 4 & 10 & 20 \end{bmatrix} = \begin{bmatrix} 1 & 1 & 1 & 1 \\ 0 & 1 & 2 & 3 \\ 0 & 2 & 5 & 9 \\ 0 & 3 & 9 & 19 \end{bmatrix}$$

Next round of G.E.

$$\begin{bmatrix} 1 & 1 & 1 & 1 \\ 0 & 1 & 2 & 3 \\ 0 & 2 & 5 & 9 \\ 0 & 3 & 9 & 19 \end{bmatrix} \xrightarrow{\begin{array}{l} R'_3 = R_3 - 2R_2 \\ R'_4 = R_4 - 3R_2 \end{array}} \begin{bmatrix} 1 & 1 & 1 & 1 \\ 0 & 1 & 2 & 3 \\ 0 & 0 & 1 & 3 \\ 0 & 0 & 3 & 10 \end{bmatrix}$$

Multiplying the corresponding elementary matrices together we have

$$\begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & -2 & 1 & 0 \\ 0 & -3 & 0 & 1 \end{bmatrix} \begin{bmatrix} 1 & 1 & 1 & 1 \\ 0 & 1 & 2 & 3 \\ 0 & 2 & 5 & 9 \\ 0 & 3 & 9 & 19 \end{bmatrix} = \begin{bmatrix} 1 & 1 & 1 & 1 \\ 0 & 1 & 2 & 3 \\ 0 & 0 & 1 & 3 \\ 0 & 0 & 3 & 10 \end{bmatrix}$$

Question

What is its inverse?

$$\begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & -2 & 1 & 0 \\ 0 & -3 & 0 & 1 \end{bmatrix}$$

Last round of G.E.

$$\begin{bmatrix} 1 & 1 & 1 & 1 \\ 0 & 1 & 2 & 3 \\ 0 & 0 & 1 & 3 \\ 0 & 0 & 3 & 10 \end{bmatrix} \xrightarrow{R'_4 = R_4 - 3R_3} \begin{bmatrix} 1 & 1 & 1 & 1 \\ 0 & 1 & 2 & 3 \\ 0 & 0 & 1 & 3 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

Now we have reached an upper triangular matrix.

$$\begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & -3 & 1 \end{bmatrix} \begin{bmatrix} 1 & 1 & 1 & 1 \\ 0 & 1 & 2 & 3 \\ 0 & 0 & 1 & 3 \\ 0 & 0 & 3 & 10 \end{bmatrix} = \begin{bmatrix} 1 & 1 & 1 & 1 \\ 0 & 1 & 2 & 3 \\ 0 & 0 & 1 & 3 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

Summary

Putting everything together:

$$\begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & -3 & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & -2 & 1 & 0 \\ 0 & -3 & 0 & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 & 0 & 0 \\ -1 & 1 & 0 & 0 \\ -1 & 0 & 1 & 0 \\ -1 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} 1 & 1 & 1 & 1 \\ 1 & 2 & 3 & 4 \\ 1 & 3 & 6 & 10 \\ 1 & 4 & 10 & 20 \end{bmatrix} = \begin{bmatrix} 1 & 1 & 1 & 1 \\ 0 & 1 & 2 & 3 \\ 0 & 0 & 1 & 3 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

third round
of G.E.
second round
of G.E.
first round
of G.E.
 A
 U

We have seen that the matrices implementing each round of Gaussian elimination are invertible.

To obtain the LU decomposition, we undo these operations to obtain an expression for A .

Undo

We undo each operation in turn, starting from the most recent.

$$\begin{array}{ccccccc}
 \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & -3 & 1 \end{bmatrix} & \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & -2 & 1 & 0 \\ 0 & -3 & 0 & 1 \end{bmatrix} & \begin{bmatrix} 1 & 0 & 0 & 0 \\ -1 & 1 & 0 & 0 \\ -1 & 0 & 1 & 0 \\ -1 & 0 & 0 & 1 \end{bmatrix} & \begin{bmatrix} 1 & 1 & 1 & 1 \\ 1 & 2 & 3 & 4 \\ 1 & 3 & 6 & 10 \\ 1 & 4 & 10 & 20 \end{bmatrix} & = & \begin{bmatrix} 1 & 1 & 1 & 1 \\ 0 & 1 & 2 & 3 \\ 0 & 0 & 1 & 3 \\ 0 & 0 & 0 & 1 \end{bmatrix} \\
 E_3 & E_2 & E_1 & A & & U
 \end{array}$$

$$\begin{array}{ccccccc}
 \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & -2 & 1 & 0 \\ 0 & -3 & 0 & 1 \end{bmatrix} & \begin{bmatrix} 1 & 0 & 0 & 0 \\ -1 & 1 & 0 & 0 \\ -1 & 0 & 1 & 0 \\ -1 & 0 & 0 & 1 \end{bmatrix} & \begin{bmatrix} 1 & 1 & 1 & 1 \\ 1 & 2 & 3 & 4 \\ 1 & 3 & 6 & 10 \\ 1 & 4 & 10 & 20 \end{bmatrix} & = & \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 3 & 1 \end{bmatrix} & \begin{bmatrix} 1 & 1 & 1 & 1 \\ 0 & 1 & 2 & 3 \\ 0 & 0 & 1 & 3 \\ 0 & 0 & 0 & 1 \end{bmatrix} \\
 E_2 & E_1 & A & & E_3^{-1} & U
 \end{array}$$

Undo

We undo each operation in turn, starting from the most recent.

$$\begin{array}{ccccc}
 \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & -2 & 1 & 0 \\ 0 & -3 & 0 & 1 \end{bmatrix} & \begin{bmatrix} 1 & 0 & 0 & 0 \\ -1 & 1 & 0 & 0 \\ -1 & 0 & 1 & 0 \\ -1 & 0 & 0 & 1 \end{bmatrix} & \begin{bmatrix} 1 & 1 & 1 & 1 \\ 1 & 2 & 3 & 4 \\ 1 & 3 & 6 & 10 \\ 1 & 4 & 10 & 20 \end{bmatrix} & = & \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 3 & 1 \end{bmatrix} \begin{bmatrix} 1 & 1 & 1 & 1 \\ 0 & 1 & 2 & 3 \\ 0 & 0 & 1 & 3 \\ 0 & 0 & 0 & 1 \end{bmatrix} \\
 E_2 & E_1 & A & & E_3^{-1} \quad U
 \end{array}$$

$$\begin{array}{ccccc}
 \begin{bmatrix} 1 & 0 & 0 & 0 \\ -1 & 1 & 0 & 0 \\ -1 & 0 & 1 & 0 \\ -1 & 0 & 0 & 1 \end{bmatrix} & \begin{bmatrix} 1 & 1 & 1 & 1 \\ 1 & 2 & 3 & 4 \\ 1 & 3 & 6 & 10 \\ 1 & 4 & 10 & 20 \end{bmatrix} & = & \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 2 & 1 & 0 \\ 0 & 3 & 0 & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 3 & 1 \end{bmatrix} \begin{bmatrix} 1 & 1 & 1 & 1 \\ 0 & 1 & 2 & 3 \\ 0 & 0 & 1 & 3 \\ 0 & 0 & 0 & 1 \end{bmatrix} \\
 E_1 & A & & E_2^{-1} \quad E_3^{-1} & U
 \end{array}$$

Undo

We undo each operation in turn, starting from the most recent.

$$\begin{bmatrix} 1 & 0 & 0 & 0 \\ -1 & 1 & 0 & 0 \\ -1 & 0 & 1 & 0 \\ -1 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} 1 & 1 & 1 & 1 \\ 1 & 2 & 3 & 4 \\ 1 & 3 & 6 & 10 \\ 1 & 4 & 10 & 20 \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 2 & 1 & 0 \\ 0 & 3 & 0 & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 3 & 1 \end{bmatrix} \begin{bmatrix} 1 & 1 & 1 & 1 \\ 0 & 1 & 2 & 3 \\ 0 & 0 & 1 & 3 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

E_1 A E_2^{-1} E_3^{-1} U

$$\begin{bmatrix} 1 & 1 & 1 & 1 \\ 1 & 2 & 3 & 4 \\ 1 & 3 & 6 & 10 \\ 1 & 4 & 10 & 20 \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 1 & 1 & 0 & 0 \\ 1 & 0 & 1 & 0 \\ 1 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 2 & 1 & 0 \\ 0 & 3 & 0 & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 3 & 1 \end{bmatrix} \begin{bmatrix} 1 & 1 & 1 & 1 \\ 0 & 1 & 2 & 3 \\ 0 & 0 & 1 & 3 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

A E_1^{-1} E_2^{-1} E_3^{-1} U

Final Result

$$\begin{array}{c}
 \begin{bmatrix} 1 & 1 & 1 & 1 \\ 1 & 2 & 3 & 4 \\ 1 & 3 & 6 & 10 \\ 1 & 4 & 10 & 20 \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 1 & 1 & 0 & 0 \\ 1 & 0 & 1 & 0 \\ 1 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 2 & 1 & 0 \\ 0 & 3 & 0 & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 3 & 1 \end{bmatrix} \begin{bmatrix} 1 & 1 & 1 & 1 \\ 0 & 1 & 2 & 3 \\ 0 & 0 & 1 & 3 \\ 0 & 0 & 0 & 1 \end{bmatrix} \\
 A \qquad E_1^{-1} \qquad E_2^{-1} \qquad E_3^{-1} \qquad U
 \end{array}$$

From the problem set, the product of lower triangular matrices is lower triangular.

Set $L = E_1^{-1} E_2^{-1} E_3^{-1}$.

$$\begin{array}{c}
 \begin{bmatrix} 1 & 1 & 1 & 1 \\ 1 & 2 & 3 & 4 \\ 1 & 3 & 6 & 10 \\ 1 & 4 & 10 & 20 \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 1 & 1 & 0 & 0 \\ 1 & 2 & 1 & 0 \\ 1 & 3 & 3 & 1 \end{bmatrix} \begin{bmatrix} 1 & 1 & 1 & 1 \\ 0 & 1 & 2 & 3 \\ 0 & 0 & 1 & 3 \\ 0 & 0 & 0 & 1 \end{bmatrix} \\
 A \qquad L \qquad U
 \end{array}$$

Big Picture

Let's review what happened.

In this Gaussian elimination we only used the operation of adding a multiple of one row **to a row below it**.

The corresponding elementary matrices are **lower triangular**.

In matrix form Gaussian elimination becomes

$$E_k \cdots E_2 E_1 A = U$$

where U is upper triangular and each E_i is lower triangular.

Big Picture

In matrix form Gaussian elimination becomes

$$E_k \cdots E_2 E_1 A = U$$

where U is upper triangular and each E_i is lower triangular.

The **inverse** of each E_i is also lower triangular.

$$A = E_1^{-1} E_2^{-1} \cdots E_k^{-1} U$$

The product of lower triangular matrices is lower triangular:

$$A = LU$$

where $L = E_1^{-1} E_2^{-1} \cdots E_k^{-1}$.

Summary

Theorem: If Gaussian elimination on the square matrix A proceeds without row swaps, then there is a factorization

$$A = LU$$

where L is lower triangular and U is upper triangular.

Row Swaps

The trouble with row swaps is simply that they are not lower triangular.

$$\begin{bmatrix} 0 & 0 & 1 \\ 0 & 1 & 0 \\ 1 & 0 & 0 \end{bmatrix}$$

We will see next time how to deal with matrices where Gaussian elimination requires row swaps.

Application

This is how **square** linear systems are actually solved in practice! Say that $A = LU$

$$\begin{bmatrix} L \end{bmatrix} \begin{bmatrix} U \end{bmatrix} \begin{bmatrix} \vec{x} \end{bmatrix} = \begin{bmatrix} \vec{b} \end{bmatrix}$$

First solve

$$\begin{bmatrix} L \end{bmatrix} \begin{bmatrix} \vec{y} \end{bmatrix} = \begin{bmatrix} \vec{b} \end{bmatrix}$$

Then

$$\begin{bmatrix} U \end{bmatrix} \begin{bmatrix} \vec{x} \end{bmatrix} = \begin{bmatrix} \vec{y} \end{bmatrix}$$

Application

Together we have

$$\begin{array}{l} L\vec{y} = \vec{b} \\ U\vec{x} = \vec{y} \end{array} \quad \longrightarrow \quad L(U\vec{x}) = \vec{b}$$

Recall that solving $U\vec{x} = \vec{y}$ for upper triangular U can be done by back substitution with about n^2 operations.

Similarly $L\vec{y} = \vec{b}$ can be solved with about n^2 operations by **forward** substitution.

Gaussian elimination, on the other hand, takes about $\frac{2}{3}n^3$ arithmetic operations.

Application

Together we have

$$\begin{array}{l} L\vec{y} = \vec{b} \\ U\vec{x} = \vec{y} \end{array} \quad \longrightarrow \quad L(U\vec{x}) = \vec{b}$$

Once we have found an LU decomposition $A = LU$, we can quickly solve $A\vec{x} = \vec{b}$ for any \vec{b} .

This is how an inverse is computed by computer (**not** with Gauss-Jordan elimination).

“inv performs an LU decomposition of the input matrix. It then uses the results to form a linear system whose solution is the matrix inverse inv(X).”

Python Experiment

Let's compare the performance in python of two ways of solving a system of linear equations.

```
>>> x = np.matmul(np.linalg.inv(A), b)
```

```
>>> x = np.linalg.solve(A, b)
```